Multi-Task Financial NLP: Stance & Sentiment Classification of FOMC Communications and News Headlines

COMP6713 — Natural Language Processing — 2026 T1

Louis Nguyen Quoc Dat Bui Nam Khanh Tran Quang Minh Phan

April 18, 2026

UNSW School of Computer Science and Engineering

Outline

Problem

Datasets

Methodology

Results

Analysis

Deployment

Summary

Problem

Two Financial NLP Classification Tasks

We tackle **two complementary 3-class classification problems** over short financial text.

Stance — FOMC Communications

Label the monetary-policy stance of a sentence from U.S. Federal Reserve meetings [1].

- **dovish** — easing bias
- **hawkish** — tightening bias
- **neutral** — no clear bias

"...the Committee decided to lower the target range for the federal funds rate." → dovish

Sentiment — Financial PhraseBank

Label the investor-facing tone of a news headline [2].

- **negative**
- **neutral**
- **positive**

"Operating profit totalled EUR 21.1 mn, up from EUR 18.6 mn a year earlier." → positive

Why It Matters

- **Central-bank stance moves markets.** A single sentence shift between hawkish and dovish language in an FOMC statement can reprice rates, equities, and FX within seconds.
- **News sentiment drives short-horizon prices.** Headline tone correlates with intraday returns and is a staple input to systematic trading and risk models.
- **Both tasks share domain vocabulary** but differ in pragmatics: stance is about *policy direction*, sentiment is about *corporate outlook*. This motivates joint modelling.
- **Labelled data is scarce and expensive.** FOMC: 2,480 sentences; FPB allagree: 2,264 sentences. Techniques that squeeze more signal out of small domain corpora are essential.
- **Practical goal:** a single deployable encoder that does both jobs competitively.

Datasets

Dataset Overview

Name	Sentences	Classes	Label distribution (test)
FOMC (Trillion Dollar Words) [1]	2,480	3 (dovish / hawkish / neutral)	dovish 130 (26.2%) hawkish 121 (24.4%) neutral 245 (49.4%)
Financial PhraseBank *allagree* [2]	2,264	3 (negative / neutral / positive)	negative 61 (13.5%) neutral 278 (61.4%) positive 114 (25.2%)

- Short texts (mostly < 40 tokens) from the financial domain.
- Both are **imbalanced**, both have **dominant neutral**, both are **small by modern NLP standards**.

Preprocessing & Splits

Splits (stratified, `random_state=42`):

- 70% train / 10% validation / 20% test
- FOMC: 1736 / 248 / 496
- FPB allagree: 1584 / 227 / 453

Text normalisation:

- Unicode NFKC, strip control chars, collapse whitespace
- Lower-case only for TF-IDF baselines
- Keep case and punctuation for BERT-family models (`max_len`=128)

Class weights for FOMC (weighted cross-entropy):

$$w_c = \frac{N}{K \cdot n_c}, \qquad w = (1.272,\ 1.365,\ 0.675)$$

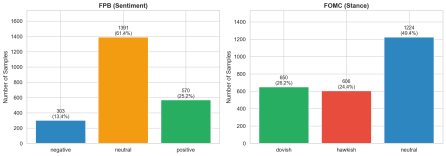

Class distributions for both tasks — neutral dominates in both, negatives and hawkish labels are the minority classes.

Methodology

Pipeline Overview

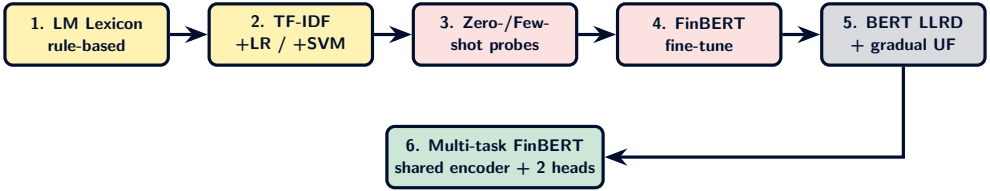

- Each stage is a **controlled upgrade** over the previous: add features, add domain prior, add fine-tuning, add training recipe, add multi-task structure.
- Every row below is evaluated on the same held-out test sets, so gains are directly comparable.

Baselines: Lexicon and TF-IDF

- **LM Lexicon (rule-based)** [3] — count Loughran–McDonald positive/negative terms per sentence, predict by sign. Domain-specific but brittle.
- **TF-IDF + Logistic Regression** — bigrams (ngram=(1,2)), `max_features`=50k, `sublinear_tf=True`; LR with default ℓ_2.
- **TF-IDF + Linear SVM** — same vectoriser; `LinearSVC`, C=1.0, `class_weight="balanced"` to offset the neutral skew.
- **TF-IDF (trigrams) + LR** — richer context: ngram=(1,3), `max_features`=80k, `min_df`=2.
- **TF-IDF + LM Lexicon (hybrid) + LR** — concatenate TF-IDF vectors with handcrafted counts of LM positive/negative terms.

These establish a *domain-aware but shallow* floor. Any neural model we propose must clearly beat them.

Pre-trained Evaluation: Zero-shot and Few-shot

Zero-shot FinBERT (native 3-way head) [4, 5]

- Use the released {positive, negative, neutral} head directly on FPB.
- For FOMC: proxy map positive $\rightarrow$ hawkish, negative $\rightarrow$ dovish, neutral $\rightarrow$ neutral. Tests whether sentiment is a usable proxy for stance.
- No parameter updates.

Few-shot linear probe (k=16 per class)

- Freeze encoder, use `[CLS]` embedding, fit logistic regression on $3 \times 16 = 48$ examples.
- Compare three frozen backbones:
  - FinBERT (`ProsusAI/finbert`)
  - BERT-base-uncased [6]
  - RoBERTa-base
- Isolates **domain pretraining value** in a tiny-data setting.

Fine-tuning: FinBERT vs. BERT-base LLRD

FinBERT single-task fine-tune

- Backbone: `ProsusAI/finbert`
- Optimiser: AdamW, lr $2e{-5}$, weight decay 0.01
- Schedule: linear warmup (10%) + linear decay
- 5 epochs, batch size 32, `max_len`=128
- Weighted cross-entropy for FOMC
- Strong domain prior, conventional recipe

BERT-base LLRD + Gradual Unfreeze [7, 8]

- Backbone: `bert-base-uncased`
- Layer-wise LR decay: head $2e{-5}$, $\times 0.9$ per depth
- Label smoothing $\varepsilon = 0.1$
- 10 epochs, batch size 32
- Gradual unfreeze: 1 layer group per epoch
- General-purpose backbone, careful recipe

The key comparison: does domain pretraining (FinBERT) beat a better *fine-tuning recipe* on a general backbone (BERT-base)? Results say: **they tie**.

Multi-task FinBERT Architecture

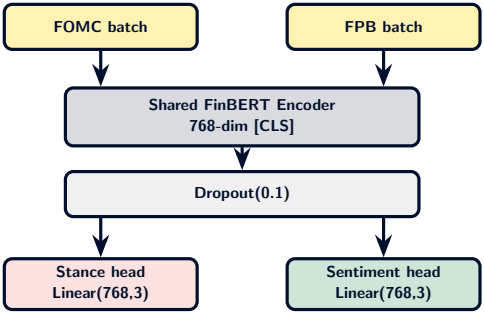

Training details [9]:

- **Alternating batches:** one FOMC batch, then one FPB batch; each updates only its own head plus the shared encoder.
- AdamW, lr $2e{-}5$, 8 epochs, batch size 32.
- Per-task cross-entropy loss (weighted for stance).
- **Best checkpoint** selected by *average* validation macro-F1 across both tasks.
- Apple M3 Max (MPS), seed 42.

Shared encoder $\Rightarrow$ stance examples regularise the sentiment head (and vice

Results

Baseline & Lexicon Results

Model	Sent. Acc	Sent. F1	Stance Acc	Stance F1
LM Lexicon (rule-based) [3]	0.6932	0.5315	0.4153	0.3885
TF-IDF + LR	0.8720	0.8232	0.6089	0.5873
TF-IDF + SVM (LinearSVC)	**0.8940**	**0.8534**	**0.6331**	**0.6061**
TF-IDF (trigrams) + LR	0.8786	0.8310	0.6109	0.5914
TF-IDF + LM Lexicon (hybrid) + LR	0.8543	0.8050	0.6109	0.5863

- **TF-IDF + LinearSVM is the strongest non-neural baseline** on both tasks.
- **Lexicon alone is weak** (FPB F1 only 0.5315, stance F1 0.3885) — sign-counting cannot handle negation or context.
- **Trigrams give only marginal lift** over bigrams (+0.5 pp accuracy) — vocabulary grows faster than signal.
- Adding LM-lexicon features to TF-IDF does *not* help — the TF-IDF vocabulary already covers those words.

Pre-trained Models: Zero-shot and Few-shot

Model	Sent. Acc	Sent. F1	Stance Acc	Stance F1
FinBERT zero-shot (native head) [4]	0.9735	0.9650	0.4980	0.4874
FinBERT few-shot (k=16 probe)	**0.9779**	**0.9670**	0.4859	0.4534
BERT-base few-shot (k=16) [6]	0.7417	0.6500	0.3851	0.3744
RoBERTa-base few-shot (k=16)	0.7682	0.6722	0.3730	0.3600

- **FinBERT zero-shot is already near-optimal on sentiment** (F1 0.9650) — its pretraining matches FPB almost exactly.
- **FinBERT dominates both probes**: +30 pp sentiment F1 over BERT/RoBERTa at k=16 — domain pretraining, not size, is what matters in low-data regimes.
- **Stance is much harder zero-shot**: the positive→hawkish / negative→dovish proxy map only reaches F1 0.4874 — stance $\neq$ sentiment.
- The probe tops sentiment but *loses* a little stance F1 vs. native head — 48 examples is too few to learn stance.

Fine-tuned Single-Task Models

Model	Sent. Acc	Sent. F1	Stance Acc	Stance F1
FinBERT (fine-tuned)	0.9669	0.9459	0.6371	0.6194
BERT-base LLRD + Gradual UF [7, 8]	**0.9757**	**0.9670**	**0.6512**	**0.6383**

- **BERT-base with LLRD + gradual unfreeze matches FinBERT on sentiment** (macro-F1 tied at 0.9670) and **exceeds it on stance** by +1.89 pp F1.
- **Takeaway:** a disciplined fine-tuning recipe on a general-purpose backbone can close the domain-pretraining gap [8].
- This is the first evidence in our study that "domain BERT" is not the only path to strong finance-domain results.

Multi-task FinBERT — Headline Result

	Sentiment	Stance
Accuracy	**0.9779**	**0.6492**
Macro-F1	**0.9666**	**0.6384**
Δ vs. single-task FinBERT	+2.07 pp F1	+1.90 pp F1

- One shared encoder, two heads, alternating batches — **improves both tasks** over single-task FinBERT [9].
- Sentiment F1 reaches 0.9666 (matching BERT LLRD's 0.9670); stance F1 reaches 0.6384 (within 0.0001 of BERT LLRD's 0.6383).
- **Free lunch:** no new data, no new parameters beyond a second 768×3 head, and training cost is comparable to single-task.

Results Summary — All 12 Systems

Model	Sent. Acc	Sent. F1	Stance Acc	Stance F1
LM Lexicon (rule-based)	0.6932	0.5315	0.4153	0.3885
TF-IDF + LR	0.8720	0.8232	0.6089	0.5873
TF-IDF + SVM (LinearSVC)	0.8940	0.8534	0.6331	0.6061
TF-IDF (trigrams) + LR	0.8786	0.8310	0.6109	0.5914
TF-IDF + LM Lexicon (hybrid)	0.8543	0.8050	0.6109	0.5863
FinBERT zero-shot (native)	0.9735	0.9650	0.4980	0.4874
FinBERT few-shot (k=16 probe)	0.9779	0.9670	0.4859	0.4534
BERT-base few-shot (k=16)	0.7417	0.6500	0.3851	0.3744
RoBERTa-base few-shot (k=16)	0.7682	0.6722	0.3730	0.3600
FinBERT (fine-tuned)	0.9669	0.9459	0.6371	0.6194
BERT-base LLRD + Gradual UF	0.9757	**0.9670**	**0.6512**	**0.6383 [BEST]**
Multi-task FinBERT	**0.9779**	**0.9666**	0.6492	**0.6384 [BEST]**

Statistical tie at the top: Multi-task FinBERT and BERT-base LLRD are **indistinguishable** within ±0.1 pp on both macro-F1s. Two very different routes (multi-task + domain backbone vs. single-task + careful recipe) reach the same ceiling.

Error Analysis — Multi-task FinBERT

Stance (FOMC test, N=496)

Class	Per-class F1
dovish	0.6174
hawkish	0.5869
neutral	**0.7109**

- Most errors sit on the **neutral vs. leaning** boundary — hedged language ("*appropriate to maintain*") is genuinely ambiguous.
- Hawkish is the weakest class (lowest count in train and most lexically diverse).

Sentiment (FPB test, N=453)

Class	Per-class F1
negative	0.9500
neutral	**0.9928**
positive	0.9569

- Only ~ 10 **errors total** across 453 test sentences — sentiment is near-saturated.
- Remaining errors are subtle *comparative* constructions ("*profit fell less than expected*").

Analysis

Key Findings

1. **Stance is a fundamentally harder task than sentiment.** Every model loses ≈ 30 pp macro-F1 going from FPB to FOMC. The gap is not about data size — the two datasets are similar — but about *label pragmatics*.
2. **Sentiment is saturated by domain pretraining alone.** FinBERT zero-shot already reaches F1 0.9650; fine-tuning adds almost nothing on FPB. Extra capacity is spent on stance.
3. **A careful fine-tuning recipe rivals domain pretraining.** BERT-base with LLRD + gradual unfreeze + label smoothing *matches* FinBERT on sentiment and *beats* it on stance [8, 7].
4. **Multi-task training is a free lunch.** One shared encoder + two heads improves single-task FinBERT by +2.07 pp sentiment F1 and +1.90 pp stance F1 [9], and ties BERT LLRD on both.
5. **Zero-shot sentiment $\neq$ stance.** The native FinBERT positive/negative head maps poorly to hawkish/dovish (F1 0.4874) — a reminder that tone and policy direction are different signals.

Domain-Pretraining Gap (k=16 probe)

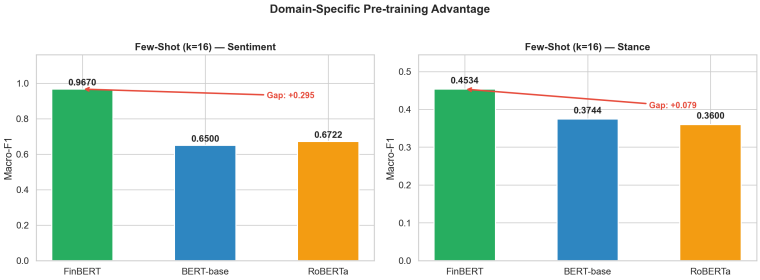

In the low-data regime, **FinBERT outperforms BERT-base and RoBERTa-base by** ≈ 30 **pp sentiment F1**. Fine-tuning closes part of this gap — but *only* when paired with a careful recipe (LLRD + gradual unfreeze).

Performance Progression Across Methods

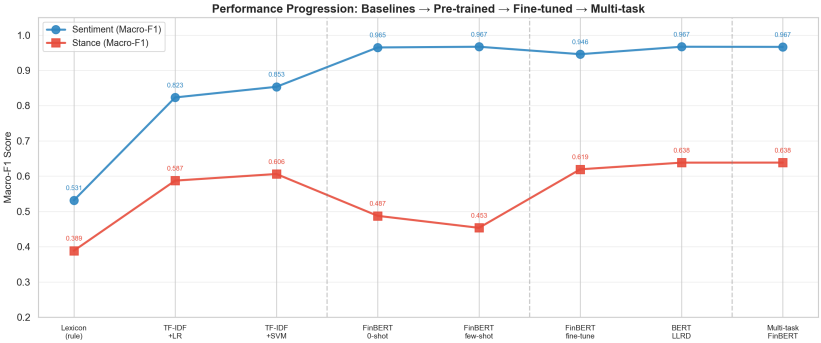

The sentiment curve **plateaus** early (FinBERT zero-shot is already near-ceiling); the stance curve keeps climbing through fine-tuning and multi-task. This is where the remaining engineering headroom lies.

Deployment

Deployment — CLI, Demo, and Hub Release

Command-line interface

- `python cli.py --text "The Committee decided to raise the target range."`
- Returns predicted stance and sentiment with class probabilities.
- Loads the multi-task FinBERT checkpoint by default; single-task checkpoints available via `--model`.

Interactive demo

- Gradio app at `http://localhost:7860`
- Text box + live predictions + per-class probability bars.
- Same checkpoint as the CLI — reproducible.

HuggingFace Hub release

Five public model repos under the `Louisnguyen/` namespace:

- `Louisnguyen/finbert-fomc-stance`
- `Louisnguyen/finbert-fpb-sentiment`
- `Louisnguyen/bert-llrd-fomc-stance`
- `Louisnguyen/bert-llrd-fpb-sentime`
- `Louisnguyen/finbert-multitask-for`

Each repo ships weights, tokenizer config, and an inference snippet.

Summary

Summary

- **Problem:** joint stance (FOMC) and sentiment (FPB) classification on small, imbalanced financial corpora.
- **Approach:** six-stage progression from LM-lexicon rules through TF-IDF, zero/few-shot probes, single-task fine-tuning (FinBERT and BERT-LLRD), and finally multi-task FinBERT with a shared encoder and two heads.
- **Headline result:** multi-task FinBERT and BERT-base LLRD *tie* for best on both tasks (sentiment F1 $\approx$ 0.967; stance F1 $\approx$ 0.638). Multi-task beats single-task FinBERT by +2.07 pp sentiment F1 / +1.90 pp stance F1.
- **Lesson:** domain pretraining and careful fine-tuning recipes are **substitutes**, not complements — either one reaches the ceiling for this data size.

Questions?

References i

[1] Agam Shah et al. **"Trillion Dollar Words: A New Financial Dataset, Task & Market Analysis".** In: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*. 2023.

[2] Pekka Malo et al. **"Good Debt or Bad Debt: Detecting Semantic Orientations in Economic Texts".** In: *Journal of the Association for Information Science and Technology* 65.4 (2014), pp. 782–796.

[3] Tim Loughran and Bill McDonald. **"When Is a Liability Not a Liability? Textual Analysis, Dictionaries, and 10-Ks".** In: *The Journal of Finance* 66.1 (2011), pp. 35–65.

[4] Dogu Araci. **"FinBERT: Financial Sentiment Analysis with Pre-trained Language Models".** In: *arXiv preprint arXiv:1908.10063* (2019).

[5] Yi Yang et al. **"FinBERT: A Pretrained Language Model for Financial Communications".** In: *arXiv preprint arXiv:2006.08097* (2020).

References ii

[6] Jacob Devlin et al. **"BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding".** In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*. 2019, pp. 4171–4186.

[7] Jeremy Howard and Sebastian Ruder. **"Universal Language Model Fine-tuning for Text Classification".** In: *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (ACL)*. 2018, pp. 328–339.

[8] Chi Sun et al. **"How to Fine-Tune BERT for Text Classification?"** In: *China National Conference on Chinese Computational Linguistics (CCL)*. 2019, pp. 194–206.

[9] Xiaodong Liu et al. **"Multi-Task Deep Neural Networks for Natural Language Understanding".** In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*. 2019, pp. 4487–4496.